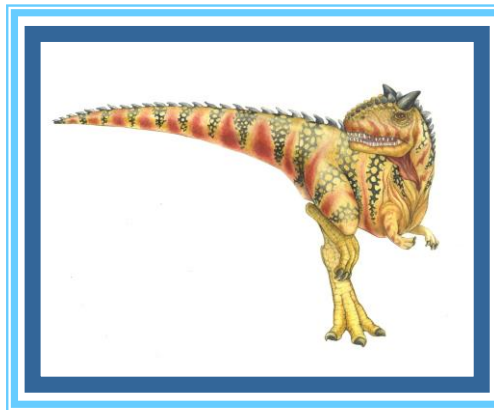


Chapter 8: Main Memory





Chapter 8: Memory Management

- ❑ Background
- ❑ Swapping
- ❑ Contiguous Memory Allocation
- ❑ Segmentation
- ❑ Paging
- ❑ Structure of the Page Table
- ❑ Example: The Intel 32 and 64-bit Architectures
- ❑ Example: ARM Architecture





Objectives

- To provide a detailed description of various ways of organizing memory hardware
- To discuss various memory-management techniques, including paging and segmentation
- To provide a detailed description of the Intel Pentium, which supports both pure segmentation and segmentation with paging





Background

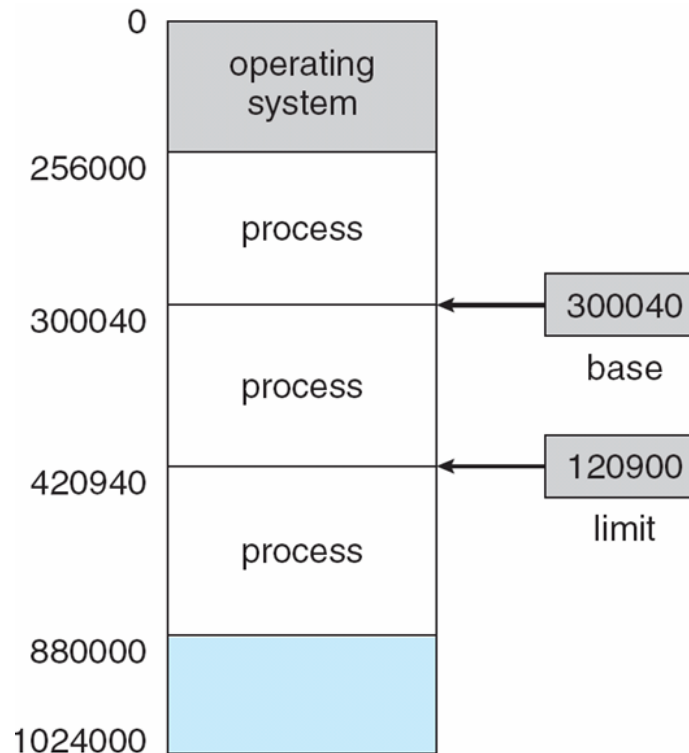
- ❑ Program must be brought (from disk) into memory and placed within a process for it to be run
- ❑ Main memory and registers are only storage CPU can access directly
- ❑ Memory unit only sees a stream of addresses + read requests, or address + data and write requests
- ❑ Register access in one CPU clock (or less)
- ❑ Main memory can take many cycles, causing a **stall**
- ❑ **Cache** sits between main memory and CPU registers
- ❑ Protection of memory required to ensure correct operation





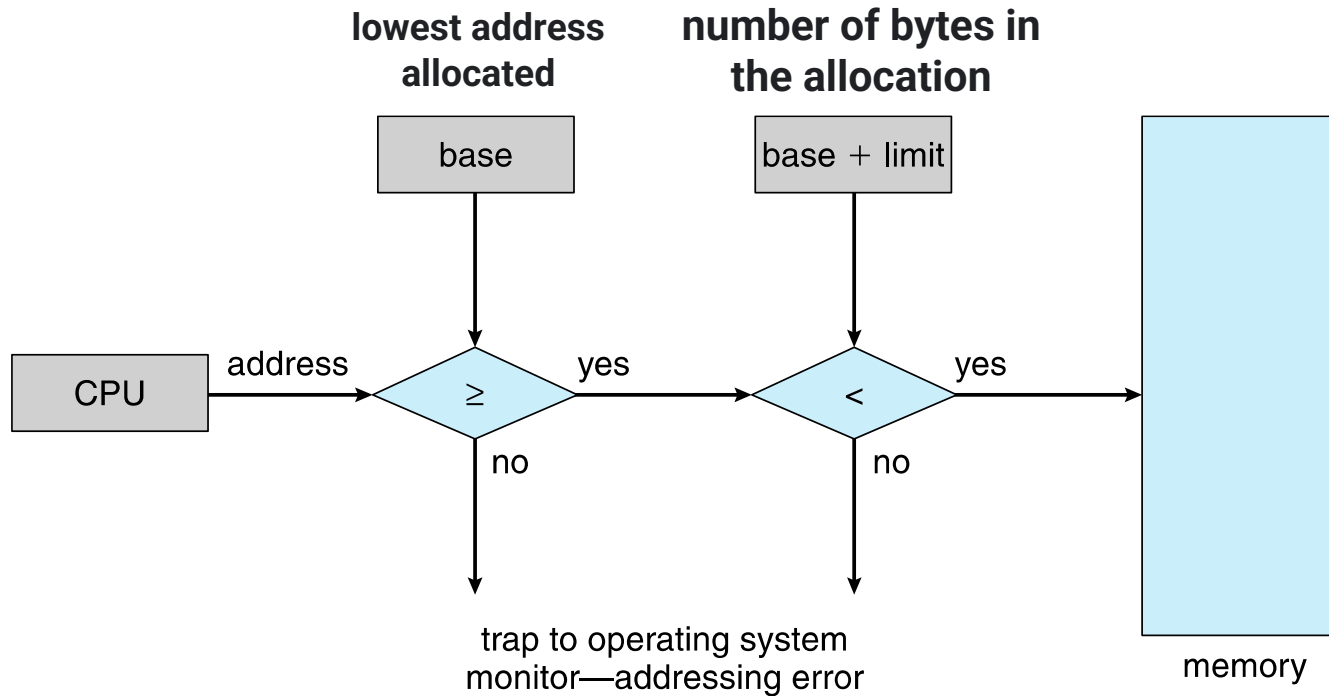
Base and Limit Registers

- A pair of **base** and **limit registers** define the logical address space
- CPU must check every memory access generated in user mode to be sure it is between base and limit for that user





Hardware Address Protection





Address Binding

- Programs on disk, ready to be brought into memory to execute form an **input queue**
 - Without support, must be loaded into address 0000
- Inconvenient to have first user process physical address always at 0000
 - How can it not be?
- Further, addresses represented in different ways at different stages of a program's life
 - Source code addresses usually symbolic
 - Compiled code addresses **bind** to relocatable addresses
 - ▶ i.e. “14 bytes from beginning of this module”
 - Linker or loader will bind relocatable addresses to absolute addresses
 - ▶ i.e. 74014
 - Each binding maps one address space to another






Linking and Loading

Loading is the process of copying an executable image into memory.

- more sophisticated loaders are able to relocate images to fit into available memory
- must readjust branch targets, load/store addresses

Linking is the process of resolving symbols between independent object files.

- suppose we define a symbol in one module, and want to use it in another
- some notation, such as `.EXTERNAL`, is used to tell assembler that a symbol is defined in another module
- linker will search symbol tables of other modules to resolve symbols and complete code generation before loading





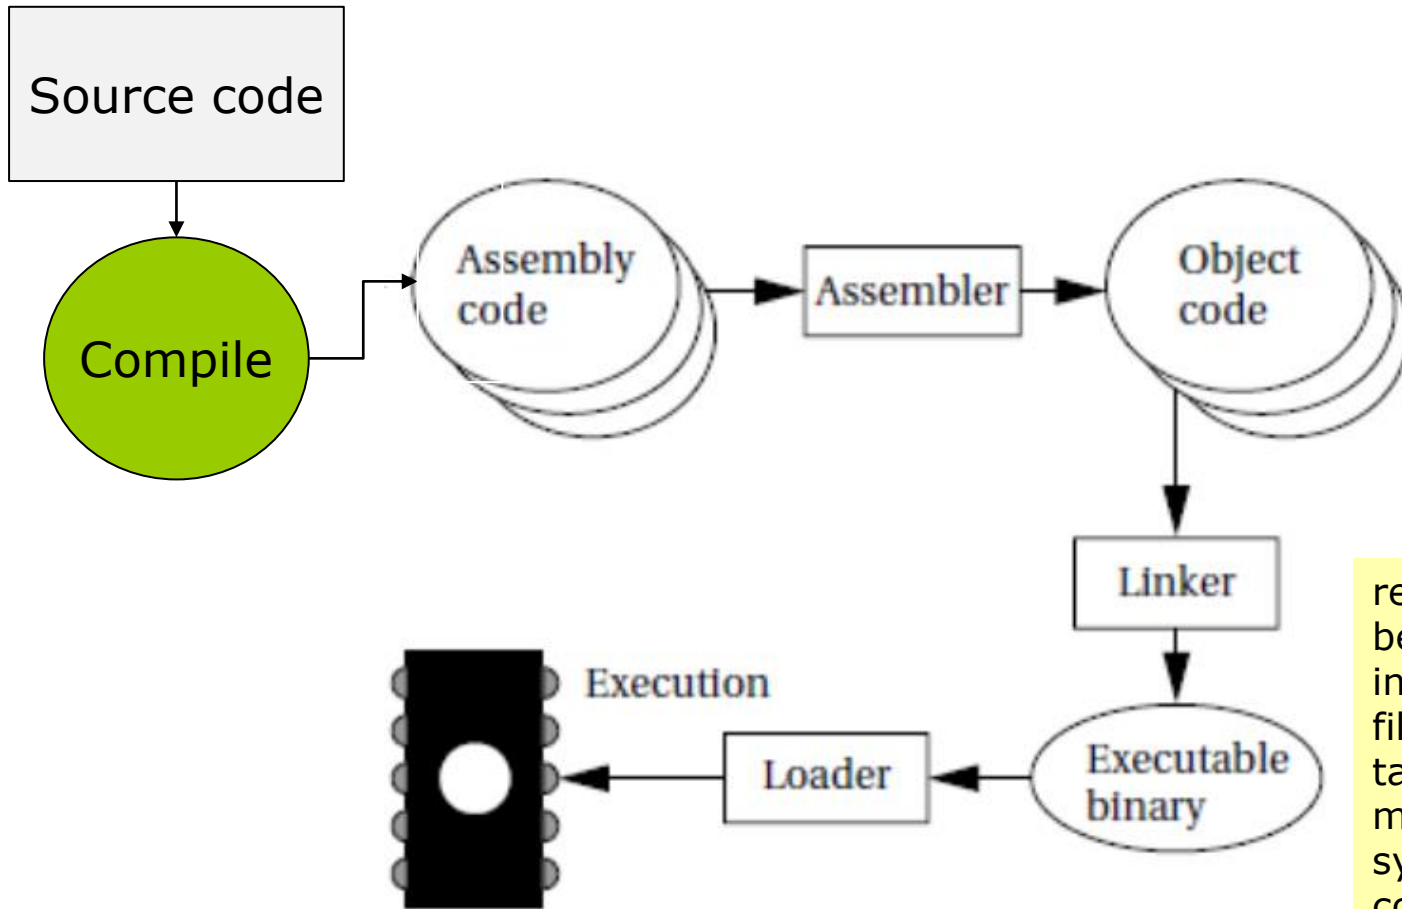
Binding of Instructions and Data to Memory

- Address binding of instructions and data to memory addresses can happen at three different stages
 - **Compile time:** If memory location known a priori, **absolute code** can be generated; must recompile code if starting location changes
 - **Load time:** Must generate **relocatable code** if memory location is not known at compile time
 - **Execution time:** Binding delayed until run time if the process can be moved during its execution from one memory segment to another
 - ▶ Need hardware support for address maps (e.g., base and limit registers)





Linking and Loading



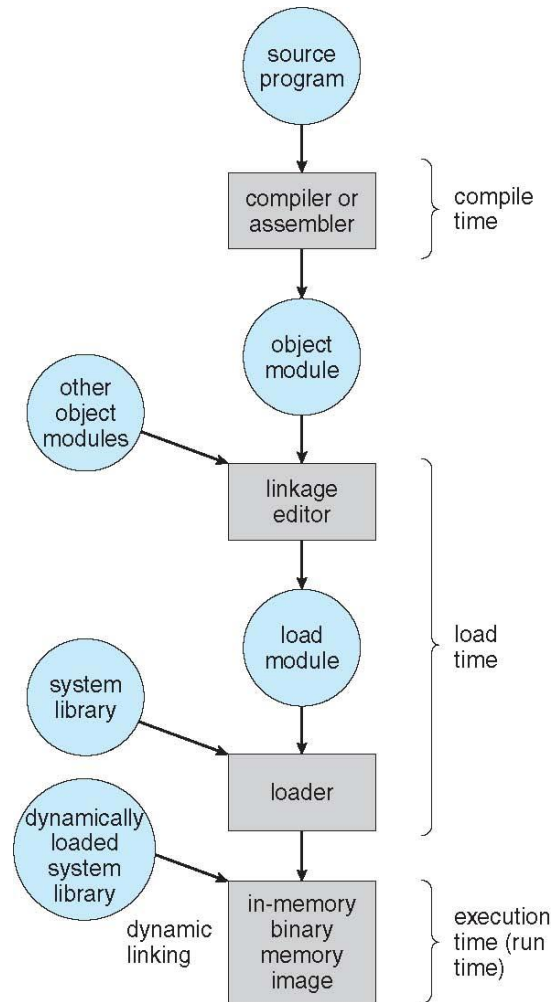
resolving symbols between independent object files: search symbol tables of other modules to resolve symbols and complete code generation before loading

copying an executable image into memory





Multistep Processing of a User Program





Logical vs. Physical Address Space

- The concept of a logical address space that is bound to a separate **physical address space** is central to proper memory management
 - **Logical address** – generated by the CPU; also referred to as **virtual address**
 - **Physical address** – address seen by the memory unit
- Logical and physical addresses are the same in compile-time and load-time address-binding schemes; logical (virtual) and physical addresses differ in execution-time address-binding scheme
- **Logical address space** is the set of all logical addresses generated by a program
- **Physical address space** is the set of all physical addresses generated by a program






Example

	Address	Instruction																Comments
1)	x30F6	1	1	1	0	0	0	1	1	1	1	1	1	1	1	0	1	$R1 \leftarrow PC - 3 = x30F4$
2)	x30F7	0	0	0	1	0	1	0	0	0	1	1	0	1	1	1	0	$R2 \leftarrow R1 + 14 = x3102$
3)	x30F8	0	0	1	1	0	1	0	1	1	1	1	1	1	0	1	1	$M[PC - 5] \leftarrow R2$ $M[x30F4] \leftarrow x3102$
4)	x30F9	0	1	0	1	0	1	0	0	1	0	1	0	0	0	0	0	$R2 \leftarrow 0$
5)	x30FA	0	0	0	1	0	1	0	0	1	0	1	0	0	1	0	1	$R2 \leftarrow R2 + 5 = 5$
6)	x30FB	0	1	1	1	0	1	0	0	0	1	0	0	1	1	1	0	$M[R1+14] \leftarrow R2$ $M[x3102] \leftarrow 5$
7)	x30FC	1	0	1	0	0	1	1	1	1	1	1	1	1	0	1	1	$R3 \leftarrow M[M[x30F4]]$ $R3 \leftarrow M[x3102]$ $R3 \leftarrow 5$

opcode






Example - LEA (Immediate)

Instruction
(1)

	15	14	13	12	11	10	9	8	7	6	5	4	3	2	1	0
LEA	1	1	1	0	Dst			PCoffset9								
	1	1	1	0	0	0	1	1	1	1	1	1	1	1	0	1


$$R1 \leftarrow PC - 3 = x30F4$$



Load Effective Address

Computes address like PC-relative (PC plus signed offset) and **stores the result into a register**.

Note: The address is stored in the register, not the contents of the memory location.




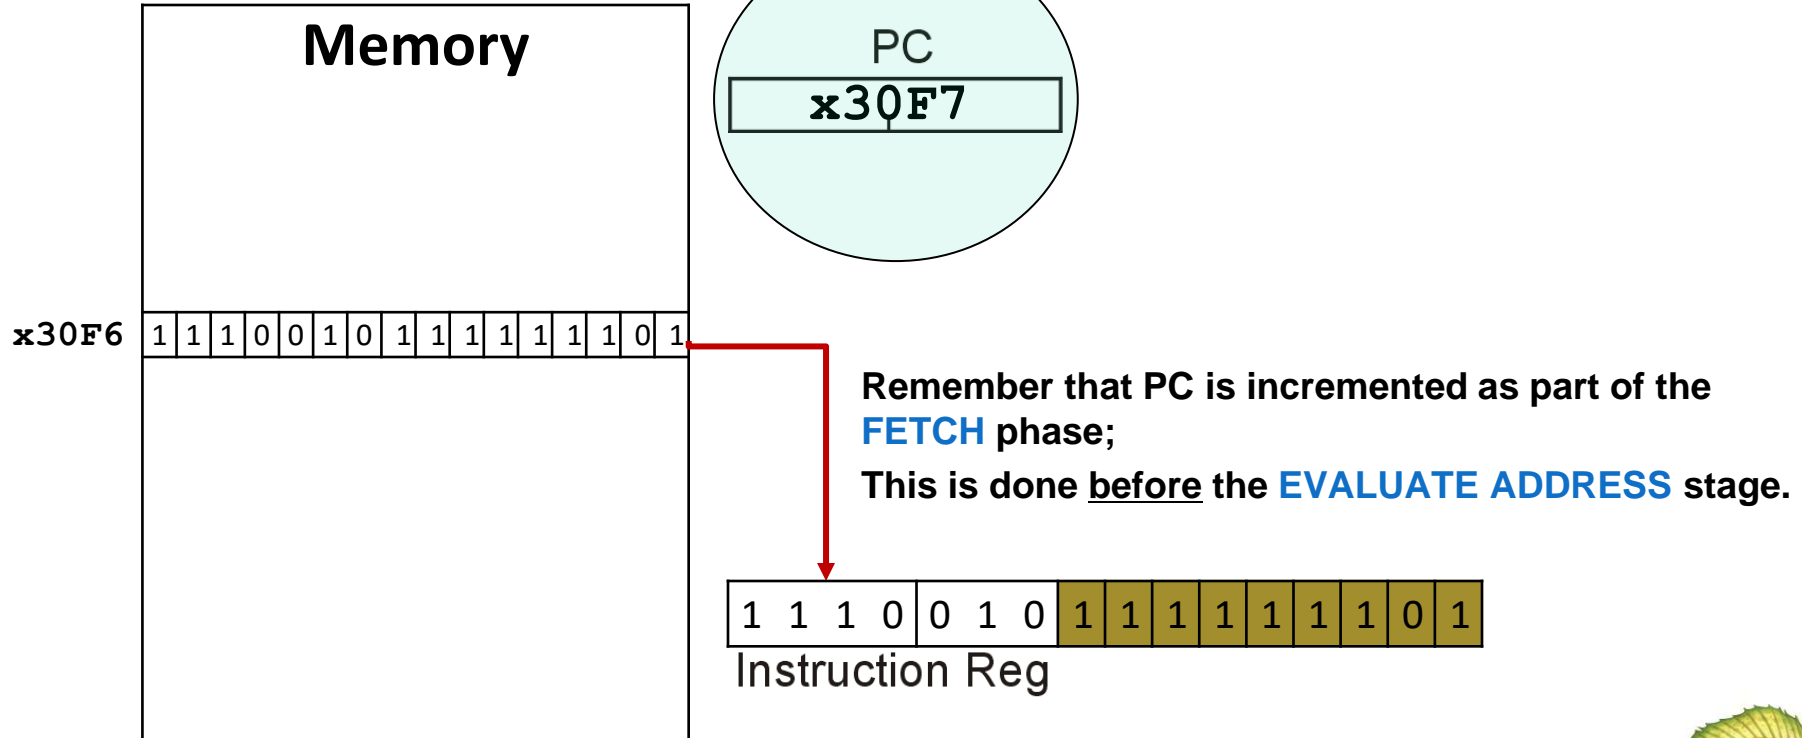


Example - LEA (Immediate)

Instruction
(1)

	15	14	13	12	11	10	9	8	7	6	5	4	3	2	1	0
LEA	1	1	1	0	Dst			PCoffset9								
	1	1	1	0	0	0	1	1	1	1	1	1	1	1	0	1

$$R1 \leftarrow PC - 3 = x30F4$$




13-15


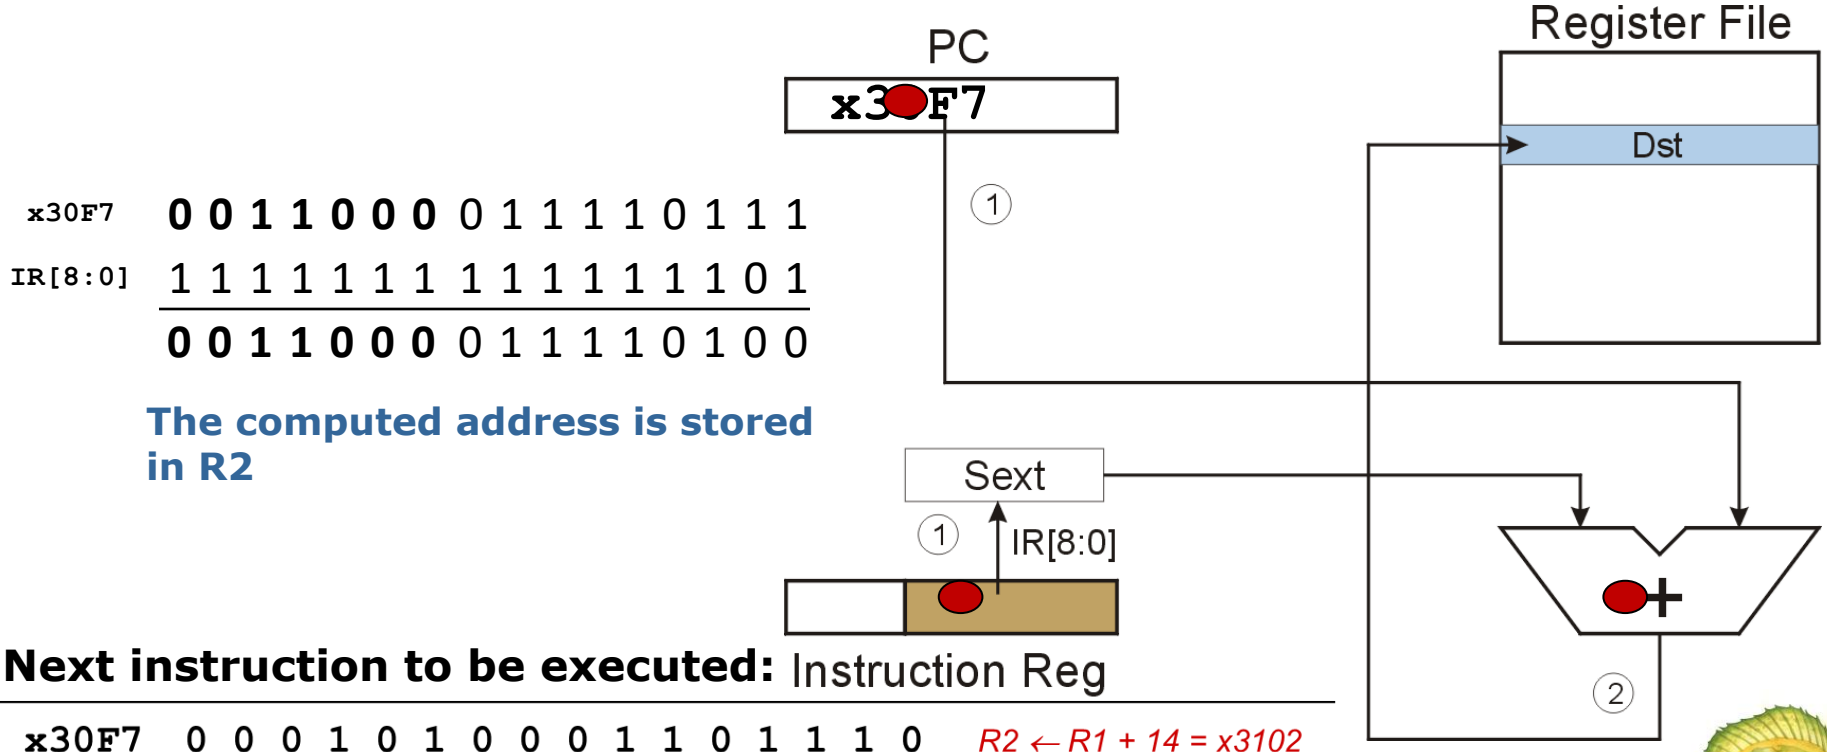


Example - LEA (Immediate)

Instruction (1)



$$R1 \leftarrow PC - 3 = x30F4$$





Memory-Management Unit (MMU)


- n Hardware device that at run time maps virtual to physical address
- n Many methods possible, covered in the rest of this chapter
- n To start, consider simple scheme where the value in the relocation register is added to every address generated by a user process at the time it is sent to memory
 - | Base register now called **relocation register**
 - | MS-DOS on Intel 80x86 used 4 relocation registers
- n The user program deals with *logical* addresses; it never sees the *real* physical addresses
 - | Execution-time binding occurs when reference is made to location in memory
 - | Logical address bound to physical addresses





Dynamic relocation using a relocation register

- n Routine is not loaded until it is called
- n Better memory-space utilization; unused routine is never loaded
- n All routines kept on disk in relocatable load format
- n Useful when large amounts of code are needed to handle infrequently occurring cases
- n No special support from the operating system is required
 - | Implemented through program design
 - | OS can help by providing libraries to implement dynamic loading





Dynamic Linking

- n **Static linking** – system libraries and program code combined by the loader into the binary program image
- n Dynamic linking –linking postponed until execution time
- n Small piece of code, **stub**, used to locate the appropriate memory-resident library routine
- n Stub replaces itself with the address of the routine, and executes the routine
- n Operating system checks if routine is in processes' memory address
 - | If not in address space, add to address space
- n Dynamic linking is particularly useful for libraries
- n System also known as **shared libraries**
- n Consider applicability to patching system libraries
 - | Versioning may be needed





Swapping

- n A process can be **swapped** temporarily out of memory to a backing store, and then brought back into memory for continued execution
 - l Total physical memory space of processes can exceed physical memory
- n **Backing store** – fast disk large enough to accommodate copies of all memory images for all users; must provide direct access to these memory images
- n **Roll out, roll in** – swapping variant used for priority-based scheduling algorithms; lower-priority process is swapped out so higher-priority process can be loaded and executed
- n Major part of swap time is transfer time; total transfer time is directly proportional to the amount of memory swapped
- n System maintains a **ready queue** of ready-to-run processes which have memory images on disk






Swapping (Cont.)

- n Does the swapped out process need to swap back in to same physical addresses?
- n Depends on address binding method
 - | Plus consider pending I/O to / from process memory space
- n Modified versions of swapping are found on many systems (i.e., UNIX, Linux, and Windows)
 - | Swapping normally disabled
 - | Started if more than threshold amount of memory allocated
 - | Disabled again once memory demand reduced below threshold





Schematic View of Swapping





Context Switch Time including Swapping

- n If next processes to be put on CPU is not in memory, need to swap out a process and swap in target process
- n Context switch time can then be very high
- n 100MB process swapping to hard disk with transfer rate of 50MB/sec
 - | Swap out time of 2000 ms
 - | Plus swap in of same sized process
 - | Total context switch swapping component time of 4000ms (4 seconds)
- n Can reduce if reduce size of memory swapped – by knowing how much memory really being used
 - | System calls to inform OS of memory use via `request_memory()` and `release_memory()`





Context Switch Time and Swapping (Cont.)

- n Other constraints as well on swapping
 - | Pending I/O – can't swap out as I/O would occur to wrong process
 - | Or always transfer I/O to kernel space, then to I/O device
 - ▶ Known as **double buffering**, adds overhead
- n Standard swapping not used in modern operating systems
 - | But modified version common
 - ▶ Swap only when free memory extremely low





Swapping on Mobile Systems

- n Not typically supported
 - | Flash memory based
 - ▶ Small amount of space
 - ▶ Limited number of write cycles
 - ▶ Poor throughput between flash memory and CPU on mobile platform
- n Instead use other methods to free memory if low
 - | iOS **asks** apps to voluntarily relinquish allocated memory
 - ▶ Read-only data thrown out and reloaded from flash if needed
 - ▶ Failure to free can result in termination
 - | Android terminates apps if low free memory, but first writes **application state** to flash for fast restart
 - | Both OSes support paging as discussed below





Contiguous Allocation

- n Main memory must support both OS and user processes
- n Limited resource, must allocate efficiently
- n Contiguous allocation is one early method
- n Main memory usually into two **partitions**:
 - | Resident operating system, usually held in low memory with interrupt vector
 - | User processes then held in high memory
 - | Each process contained in single contiguous section of memory





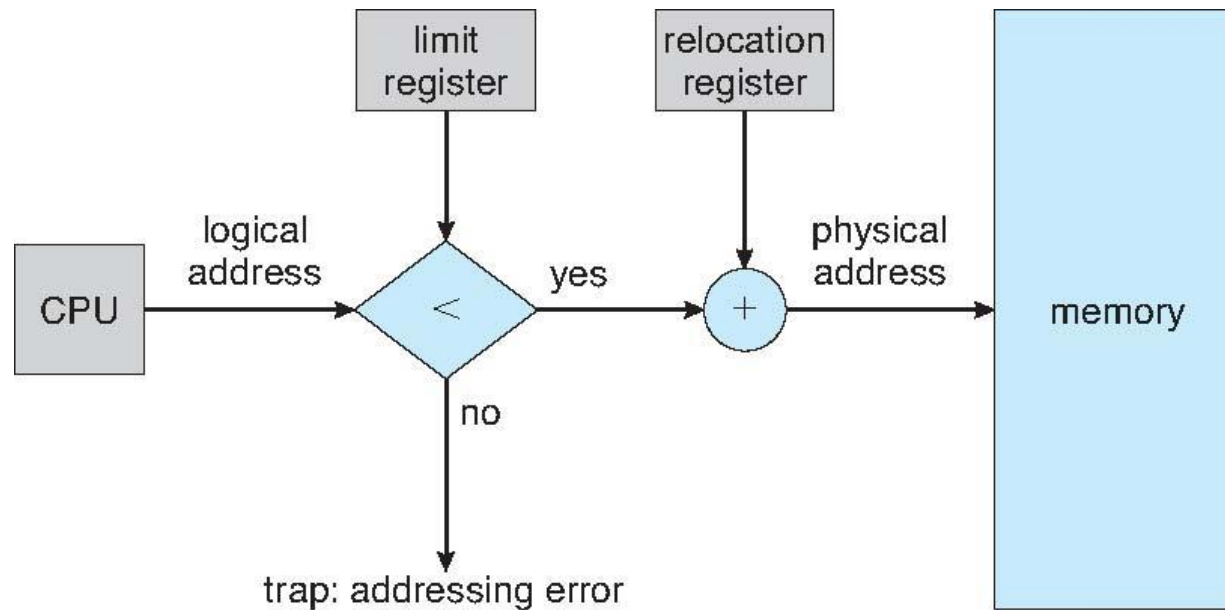
Contiguous Allocation (Cont.)

- n Relocation registers used to protect user processes from each other, and from changing operating-system code and data
 - | Base register contains value of smallest physical address
 - | Limit register contains range of logical addresses – each logical address must be less than the limit register
 - | MMU maps logical address *dynamically*
 - | Can then allow actions such as kernel code being **transient** and kernel changing size





Hardware Support for Relocation and Limit Registers






Multiple-partition allocation

n Multiple-partition allocation

- | Degree of multiprogramming limited by number of partitions
- | **Variable-partition** sizes for efficiency (sized to a given process' needs)
- | **Hole** – block of available memory; holes of various size are scattered throughout memory
- | When a process arrives, it is allocated memory from a hole large enough to accommodate it
- | Process exiting frees its partition, adjacent free partitions combined
- | Operating system maintains information about:
a) allocated partitions b) free partitions (hole)





Dynamic Storage-Allocation Problem

How to satisfy a request of size n from a list of free holes?

- n **First-fit**: Allocate the **first** hole that is big enough
- n **Best-fit**: Allocate the **smallest** hole that is big enough; must search entire list, unless ordered by size
 - | Produces the smallest leftover hole
- n **Worst-fit**: Allocate the **largest** hole; must also search entire list
 - | Produces the largest leftover hole

First-fit and best-fit better than worst-fit in terms of speed and storage utilization





Fragmentation

- n **External Fragmentation** – total memory space exists to satisfy a request, but it is not contiguous
- n **Internal Fragmentation** – allocated memory may be slightly larger than requested memory; this size difference is memory internal to a partition, but not being used
- n First fit analysis reveals that given N blocks allocated, $0.5 N$ blocks lost to fragmentation
 - | $1/3$ may be unusable -> **50-percent rule**





Fragmentation (Cont.)

- n Reduce external fragmentation by **compaction**
 - | Shuffle memory contents to place all free memory together in one large block
 - | Compaction is possible *only* if relocation is dynamic, and is done at execution time
 - | I/O problem
 - ▶ Latch job in memory while it is involved in I/O
 - ▶ Do I/O only into OS buffers
- n Now consider that backing store has same fragmentation problems






Segmentation

- n Memory-management scheme that supports user view of memory
- n A program is a collection of segments
 - | A segment is a logical unit such as:
 - main program
 - procedure
 - function
 - method
 - object
 - local variables, global variables
 - common block
 - stack
 - symbol table
 - arrays






User's View of a Program






Logical View of Segmentation



user space



physical memory space





Segmentation Architecture

- n Logical address consists of a two tuple:
 <segment-number, offset>,
- n **Segment table** – maps two-dimensional physical addresses; each table entry has:
 - | **base** – contains the starting physical address where the segments reside in memory
 - | **limit** – specifies the length of the segment
- n **Segment-table base register (STBR)** points to the segment table's location in memory
- n **Segment-table length register (STLR)** indicates number of segments used by a program;
 segment number **s** is legal if **s** < **STLR**





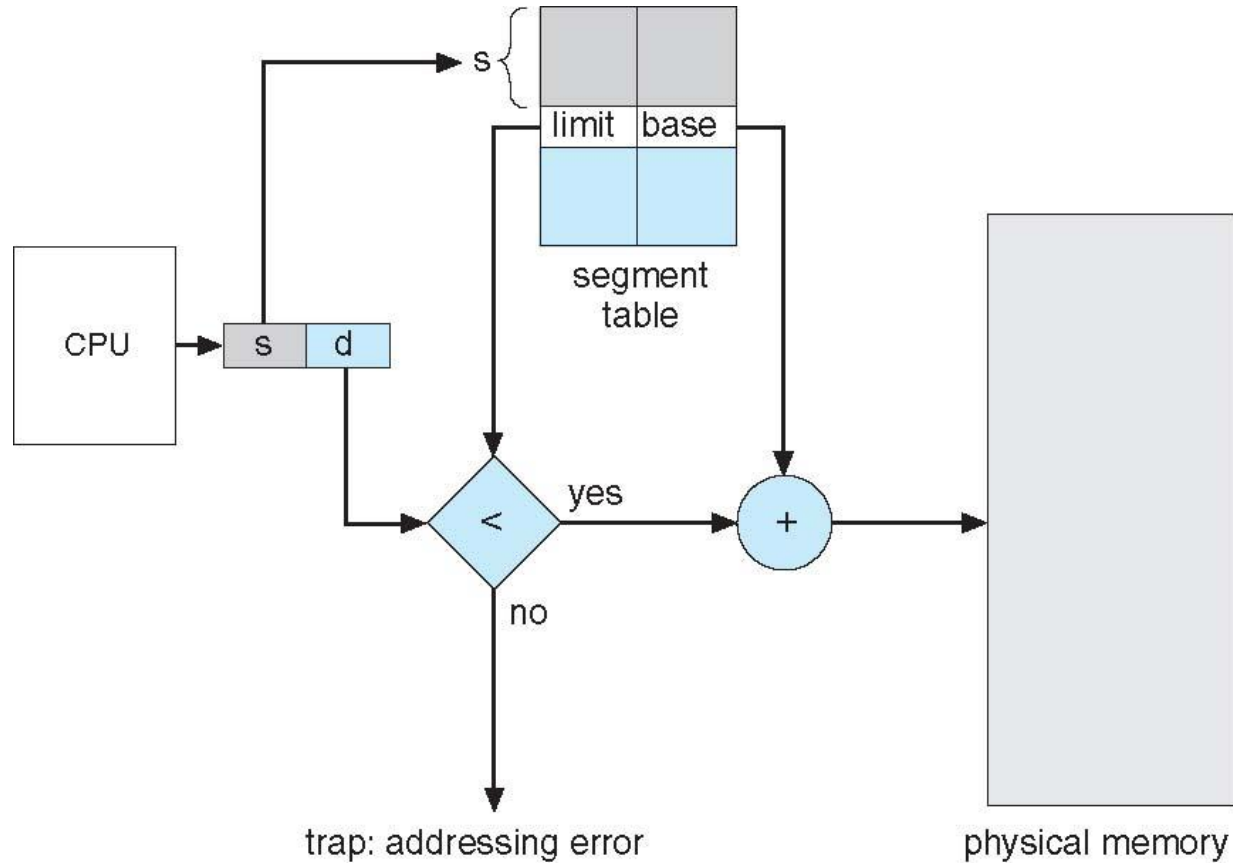
Segmentation Architecture (Cont.)

- n Protection
 - | With each entry in segment table associate:
 - ▶ validation bit = 0 \Rightarrow illegal segment
 - ▶ read/write/execute privileges
- n Protection bits associated with segments; code sharing occurs at segment level
- n Since segments vary in length, memory allocation is a dynamic storage-allocation problem
- n A segmentation example is shown in the following diagram





Segmentation Hardware





Paging


- n Physical address space of a process can be noncontiguous; process is allocated physical memory whenever the latter is available
 - | Avoids external fragmentation
 - | Avoids problem of varying sized memory chunks
- n Divide physical memory into fixed-sized blocks called **frames**
 - | Size is power of 2, between 512 bytes and 16 Mbytes
- n Divide logical memory into blocks of same size called **pages**
- n Keep track of all free frames
- n To run a program of size ***N*** pages, need to find ***N*** free frames and load program
- n Set up a **page table** to translate logical to physical addresses
- n Backing store likewise split into pages
- n Still have Internal fragmentation





Address Translation Scheme

- Address generated by CPU is divided into:
 - **Page number** (p) – used as an index into a **page table** which contains base address of each page in physical memory
 - **Page offset** (d) – combined with base address to define the physical memory address that is sent to the memory unit

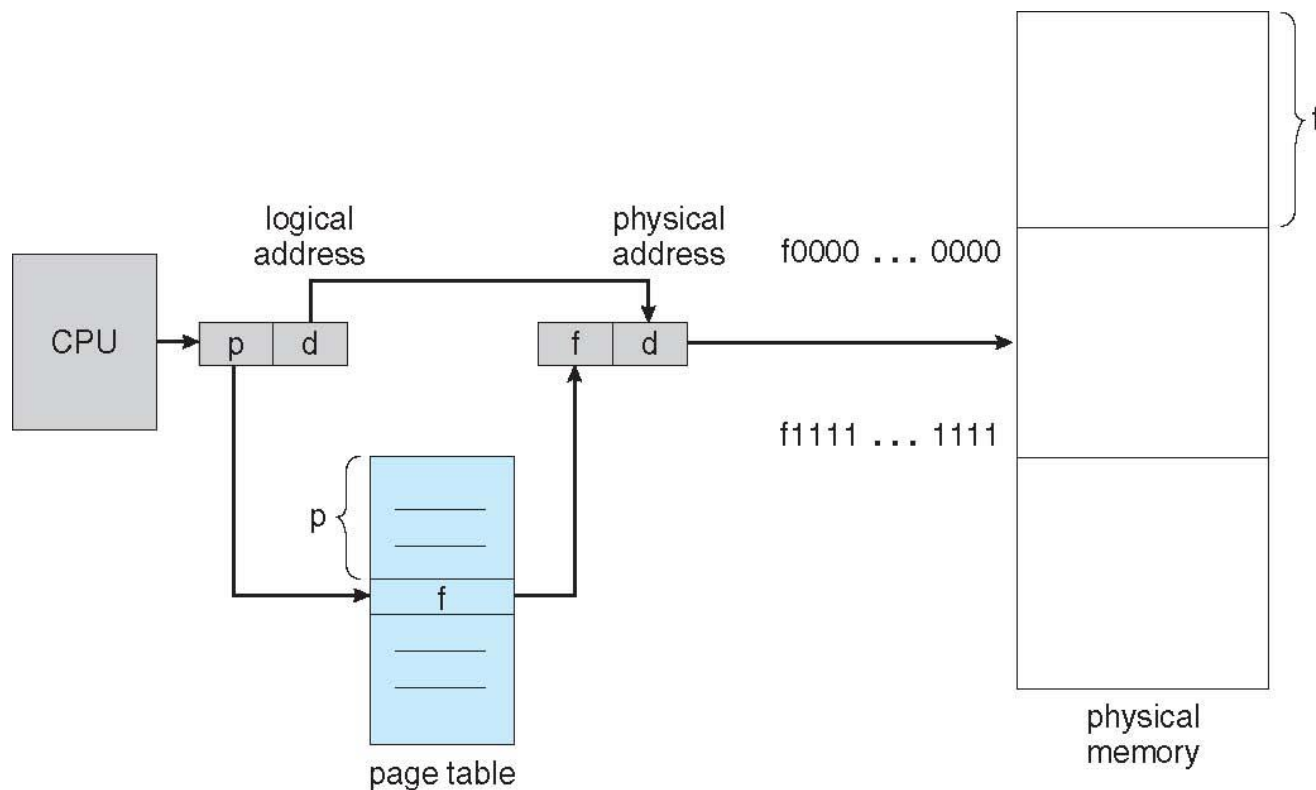


- For given logical address space 2^m and page size 2^n



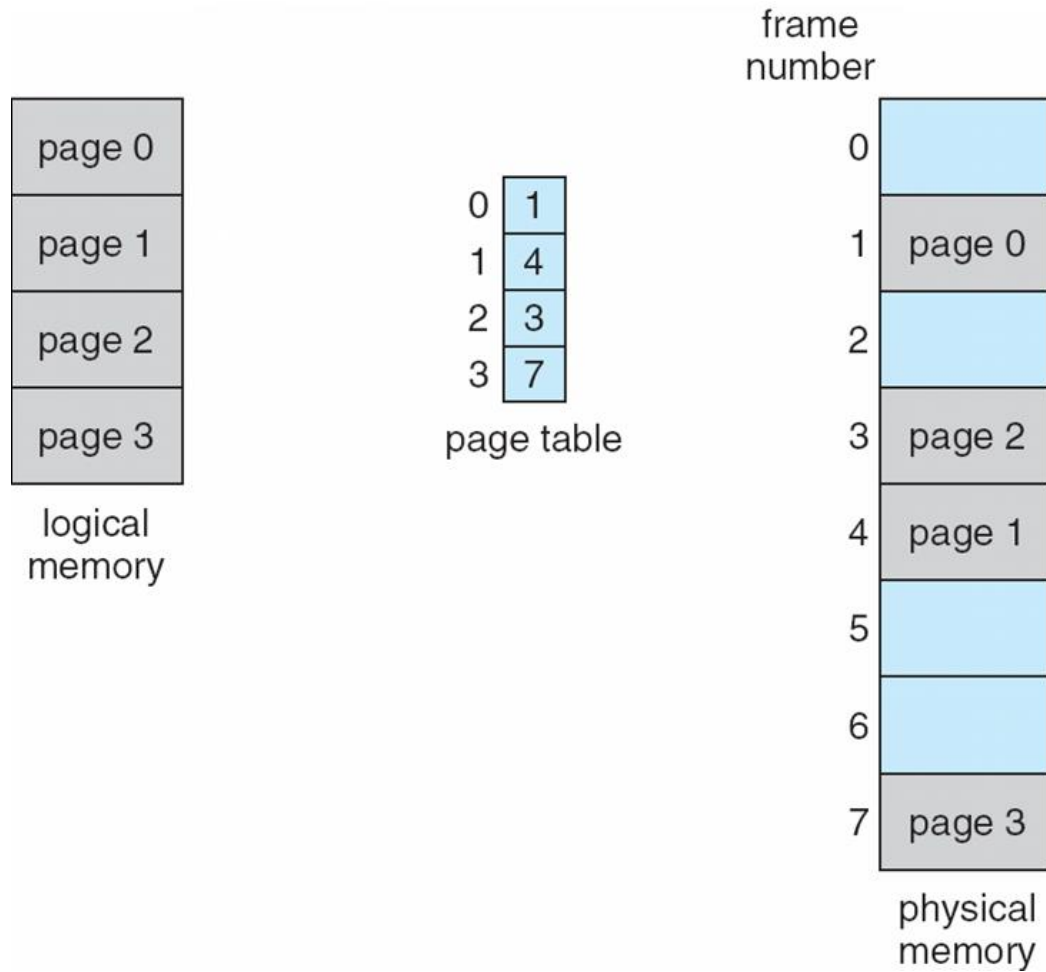


Paging Hardware






Paging Model of Logical and Physical Memory





Paging Example



$n=2$ and $m=4$ 32-byte memory and 4-byte pages






Paging (Cont.)

- Calculating internal fragmentation
 - Page size = 2,048 bytes
 - Process size = 72,766 bytes
 - 35 pages + 1,086 bytes
 - Internal fragmentation of $2,048 - 1,086 = 962$ bytes
 - Worst case fragmentation = 1 frame – 1 byte
 - On average fragmentation = $1 / 2$ frame size
 - So small frame sizes desirable?
 - But each page table entry takes memory to track
 - Page sizes growing over time
 - ▶ Solaris supports two page sizes – 8 KB and 4 MB
- Process view and physical memory now very different
- By implementation process can only access its own memory





Free Frames



Before allocation

After allocation





Implementation of Page Table

- Page table is kept in main memory
- **Page-table base register (PTBR)** points to the page table
- **Page-table length register (PTLR)** indicates size of the page table
- In this scheme every data/instruction access requires two memory accesses
 - One for the page table and one for the data / instruction
- The two memory access problem can be solved by the use of a special fast-lookup hardware cache called **associative memory** or **translation look-aside buffers (TLBs)**





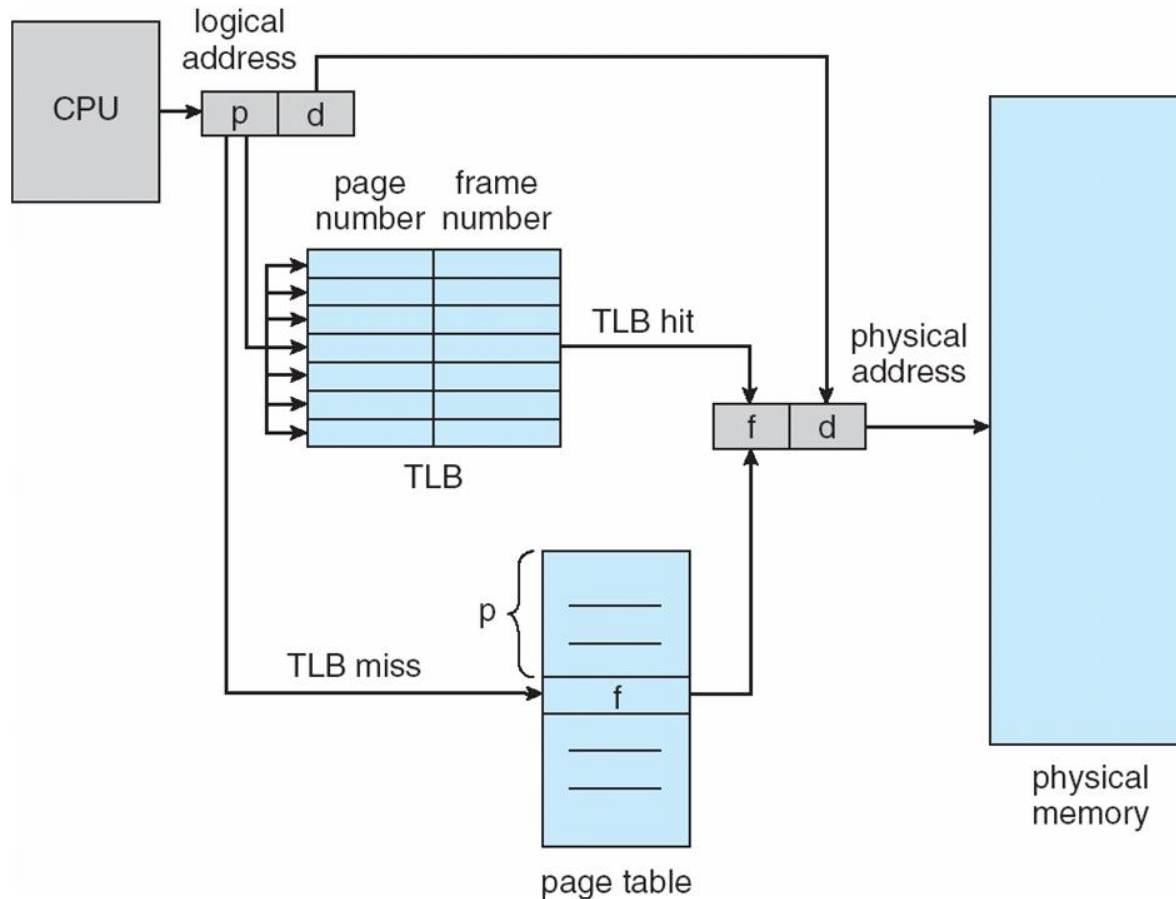
Implementation of Page Table (Cont.)

- Some TLBs store **address-space identifiers (ASIDs)** in each TLB entry – uniquely identifies each process to provide address-space protection for that process
 - Otherwise need to flush at every context switch
- TLBs typically small (64 to 1,024 entries)
- On a TLB miss, value is loaded into the TLB for faster access next time
 - Replacement policies must be considered
 - Some entries can be **wired down** for permanent fast access





Paging Hardware With TLB






Memory Protection

- Memory protection implemented by associating protection bit with each frame to indicate if read-only or read-write access is allowed
 - Can also add more bits to indicate page execute-only, and so on
- **Valid-invalid** bit attached to each entry in the page table:
 - “valid” indicates that the associated page is in the process’ logical address space, and is thus a legal page
 - “invalid” indicates that the page is not in the process’ logical address space
 - Or use **page-table length register (PTLR)**
- Any violations result in a trap to the kernel





Valid (v) or Invalid (i) Bit In A Page Table





Shared Pages

□ Shared code

- One copy of read-only (**reentrant**) code shared among processes (i.e., text editors, compilers, window systems)
- Similar to multiple threads sharing the same process space
- Also useful for interprocess communication if sharing of read-write pages is allowed

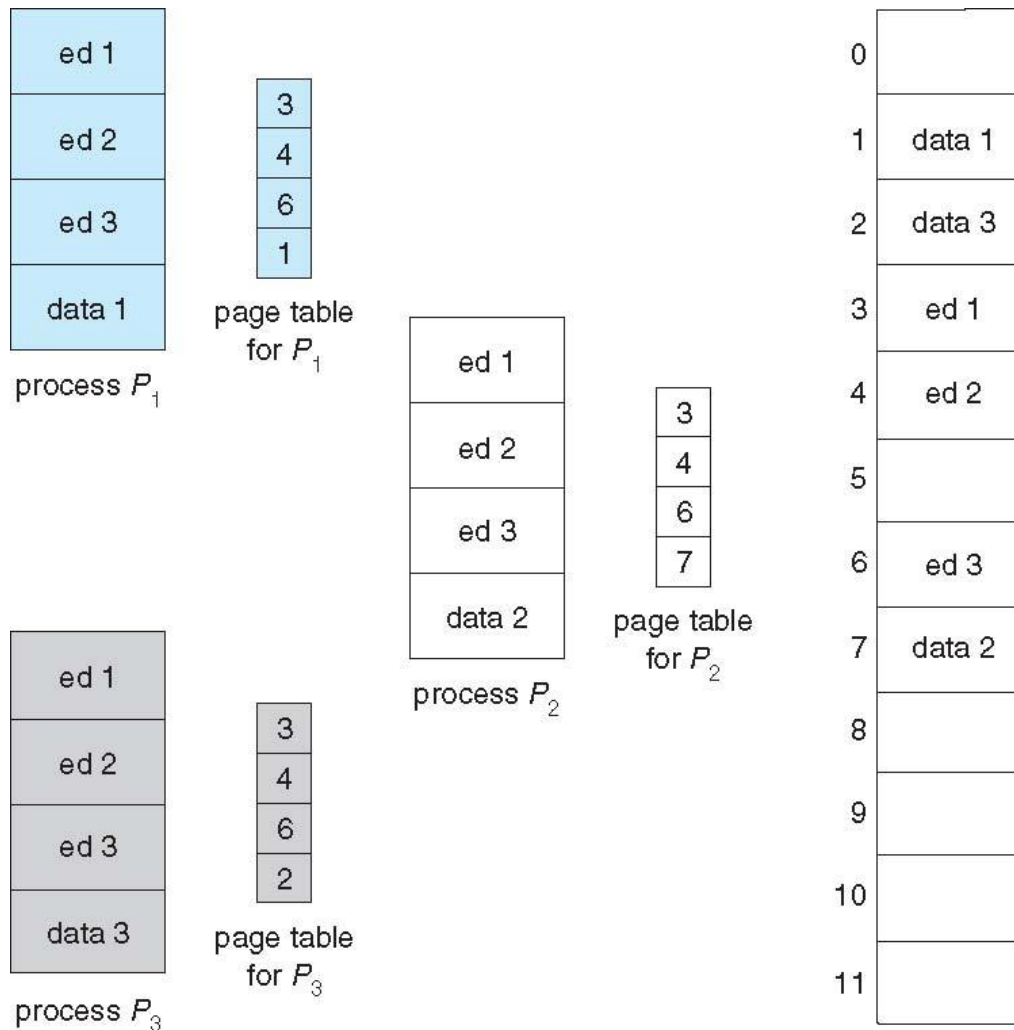
□ Private code and data

- Each process keeps a separate copy of the code and data
- The pages for the private code and data can appear anywhere in the logical address space





Shared Pages Example





Structure of the Page Table

- ❑ Memory structures for paging can get huge using straightforward methods
 - ❑ Consider a 32-bit logical address space as on modern computers
 - ❑ Page size of 4 KB (2^{12})
 - ❑ Page table would have 1 million entries ($2^{32} / 2^{12}$)
 - ❑ If each entry is 4 bytes -> 4 MB of physical address space / memory for page table alone
 - ▶ That amount of memory used to cost a lot
 - ▶ Don't want to allocate that contiguously in main memory
- ❑ Hierarchical Paging
- ❑ Hashed Page Tables
- ❑ Inverted Page Tables



End of Chapter 8

