


Chapter 9: Virtual Memory





Chapter 9: Virtual Memory

- ❑ Background
- ❑ Demand Paging
- ❑ Copy-on-Write
- ❑ Page Replacement
- ❑ Allocation of Frames
- ❑ Thrashing
- ❑ Memory-Mapped Files
- ❑ Allocating Kernel Memory
- ❑ Other Considerations
- ❑ Operating-System Examples





Objectives

- To describe the benefits of a virtual memory system
- To explain the concepts of demand paging, page-replacement algorithms, and allocation of page frames
- To discuss the principle of the working-set model
- To examine the relationship between shared memory and memory-mapped files
- To explore how kernel memory is managed





Background

- ❑ Code needs to be in memory to execute, but entire program rarely used
 - ❑ Error code, unusual routines, large data structures
- ❑ Entire program code not needed at same time
- ❑ Consider ability to execute partially-loaded program
 - ❑ Program no longer constrained by limits of physical memory
 - ❑ Each program takes less memory while running -> more programs run at the same time
 - ▶ Increased CPU utilization and throughput with no increase in response time or turnaround time
 - ❑ Less I/O needed to load or swap programs into memory -> each user program runs faster





Background (Cont.)

- **Virtual memory** – separation of user logical memory from physical memory
 - Only part of the program needs to be in memory for execution
 - Logical address space can therefore be much larger than physical address space
 - Allows address spaces to be shared by several processes
 - Allows for more efficient process creation
 - More programs running concurrently
 - Less I/O needed to load or swap processes






Background (Cont.)

- **Virtual address space** – logical view of how process is stored in memory
 - Usually start at address 0, contiguous addresses until end of space
 - Meanwhile, physical memory organized in page frames
 - MMU must map logical to physical
- Virtual memory can be implemented via:
 - Demand paging
 - Demand segmentation






Virtual Memory That is Larger Than Physical Memory





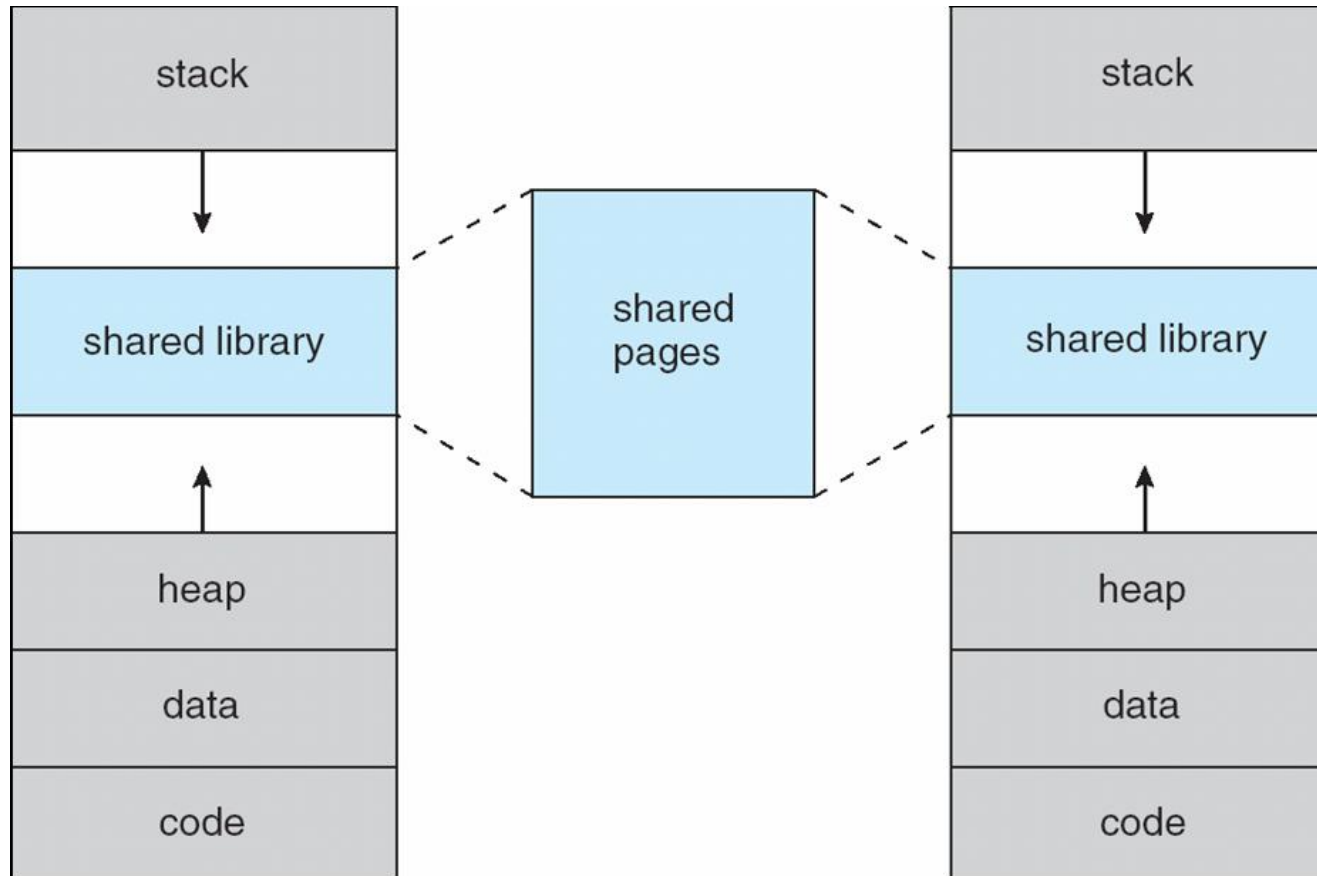
Virtual-address Space

- Usually design logical address space for stack to start at Max logical address and grow “down” while heap grows “up”
 - Maximizes address space use
 - Unused address space between the two is hole
 - ▶ No physical memory needed until heap or stack grows to a given new page
- Enables **sparse** address spaces with holes left for growth, dynamically linked libraries, etc
- System libraries shared via mapping into virtual address space
- Shared memory by mapping pages read-write into virtual address space
- Pages can be shared during `fork()`, speeding process creation






Shared Library Using Virtual Memory





Demand Paging

- ❑ Could bring entire process into memory at load time
- ❑ Or bring a page into memory only when it is needed
 - ❑ Less I/O needed, no unnecessary I/O
 - ❑ Less memory needed
 - ❑ Faster response
 - ❑ More users
- ❑ Similar to paging system with swapping (diagram on right)
- ❑ Page is needed \Rightarrow reference to it
 - ❑ invalid reference \Rightarrow abort
 - ❑ not-in-memory \Rightarrow bring to memory
- ❑ **Lazy swapper** – never swaps a page into memory unless page will be needed
 - ❑ Swapper that deals with pages is a **pager**





Basic Concepts

- With swapping, pager guesses which pages will be used before swapping out again
- Instead, pager brings in only those pages into memory
- How to determine that set of pages?
 - Need new MMU functionality to implement demand paging
- If pages needed are already **memory resident**
 - No difference from non demand-paging
- If page needed and not memory resident
 - Need to detect and load the page into memory from storage
 - ▶ Without changing program behavior
 - ▶ Without programmer needing to change code





Valid-Invalid Bit

- With each page table entry a valid–invalid bit is associated (**v** \Rightarrow in-memory – **memory resident**, **i** \Rightarrow not-in-memory)
- Initially valid–invalid bit is set to **i** on all entries
- Example of a page table snapshot:

Frame #	valid-invalid bit
	v
	v
	v
	i
...	
	i
	i

page table

- During MMU address translation, if valid–invalid bit in page table entry is **i** \Rightarrow page fault





Page Table When Some Pages Are Not in Main Memory

0	A
1	B
2	C
3	D
4	E
5	F
6	G
7	H


logical
memory

valid-invalid bit		
frame		
0	4	v
1		i
2	6	v
3		i
4		i
5	9	v
6		i
7		i

page table

0
1
2
3
4
5
6
7
8
9
10
11
12
13
14
15

physical memory





Page Fault

- If there is a reference to a page, first reference to that page will trap to operating system:

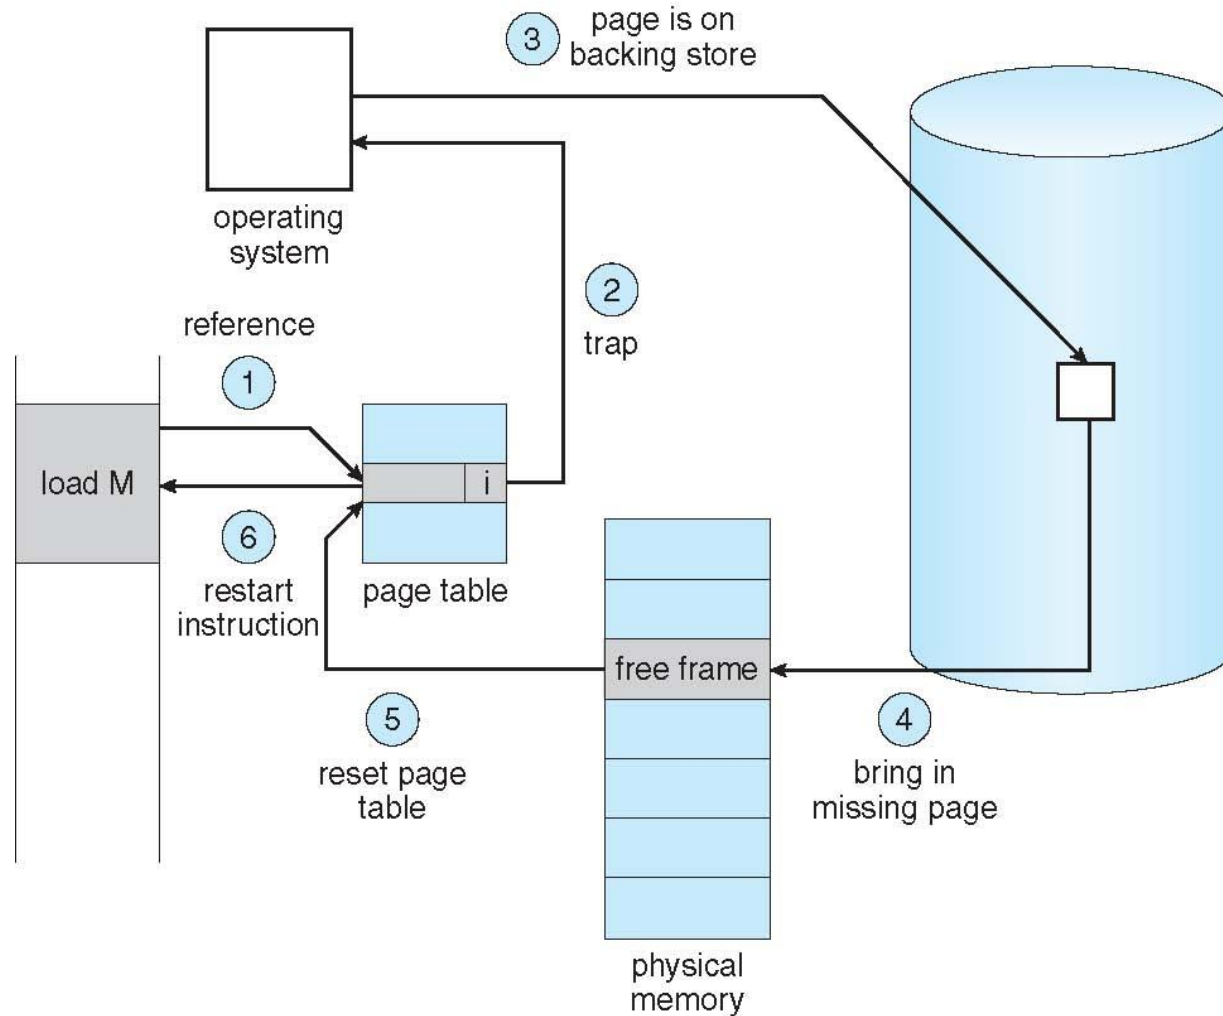
page fault

1. Operating system looks at another table to decide:
 - Invalid reference \Rightarrow abort
 - Just not in memory
2. Find free frame
3. Swap page into frame via scheduled disk operation
4. Reset tables to indicate page now in memory
Set validation bit = **v**
5. Restart the instruction that caused the page fault





Steps in Handling a Page Fault





Aspects of Demand Paging

- Extreme case – start process with *no* pages in memory
 - OS sets instruction pointer to first instruction of process, non-memory-resident -> page fault
 - And for every other process pages on first access
 - **Pure demand paging**
- Actually, a given instruction could access multiple pages -> multiple page faults
 - Consider fetch and decode of instruction which adds 2 numbers from memory and stores result back to memory
 - Pain decreased because of **locality of reference**
- Hardware support needed for demand paging
 - Page table with valid / invalid bit
 - Secondary memory (swap device with **swap space**)
 - Instruction restart





Performance of Demand Paging

□ Stages in Demand Paging (worse case)

1. Trap to the operating system
2. Save the user registers and process state
3. Determine that the interrupt was a page fault
4. Check that the page reference was legal and determine the location of the page on the disk
5. Issue a read from the disk to a free frame:
 1. Wait in a queue for this device until the read request is serviced
 2. Wait for the device seek and/or latency time
 3. Begin the transfer of the page to a free frame
6. While waiting, allocate the CPU to some other user
7. Receive an interrupt from the disk I/O subsystem (I/O completed)
8. Save the registers and process state for the other user
9. Determine that the interrupt was from the disk
10. Correct the page table and other tables to show page is now in memory
11. Wait for the CPU to be allocated to this process again
12. Restore the user registers, process state, and new page table, and then resume the interrupted instruction





Performance of Demand Paging (Cont.)

- Three major activities
 - Service the interrupt – careful coding means just several hundred instructions needed
 - Read the page – lots of time
 - Restart the process – again just a small amount of time
- Page Fault Rate $0 \leq p \leq 1$
 - if $p = 0$ no page faults
 - if $p = 1$, every reference is a fault
- Effective Access Time (EAT)

$$\begin{aligned} \text{EAT} = & (1 - p) \times \text{memory access} \\ & + p (\text{page fault overhead} \\ & \quad + \text{swap page out} \\ & \quad + \text{swap page in}) \end{aligned}$$





Page and Frame Replacement Algorithms


- **Frame-allocation algorithm** determines
 - How many frames to give each process
 - Which frames to replace
- **Page-replacement algorithm**
 - Want lowest page-fault rate on both first access and re-access
- Evaluate algorithm by running it on a particular string of memory references (reference string) and computing the number of page faults on that string
 - String is just page numbers, not full addresses
 - Repeated access to the same page does not cause a page fault
 - Results depend on number of frames available
- In all our examples, the **reference string** of referenced page numbers is

7,0,1,2,0,3,0,4,2,3,0,3,0,3,2,1,2,0,1,7,0,1





Graph of Page Faults Versus The Number of Frames





First-In-First-Out (FIFO) Algorithm

- Reference string: **7,0,1,2,0,3,0,4,2,3,0,3,2,1,2,0,1,7,0,1**
- 3 frames (3 pages can be in memory at a time per process)

reference string

7 0 1 2 0 3 0 4 2 3 0 3 2 1 2 0 1 7 0 1

7	7	7	2	2	2	4	4	4	0	0	0	0	0	0	7	7	7
	0	0	0	3	3	3	2	2	2	1	1	1	1	1	1	0	0
		1	1	1	0	0	0	0	3	3	3	2	2	2	2	2	1

page frames

15 page faults

- Can vary by reference string: consider 1,2,3,4,1,2,5,1,2,3,4,5
 - Adding more frames can cause more page faults!
 - **Belady's Anomaly**





reference string

7 0 1 2 0 3 0 4 2 3 0 3 2 1 2 0 1 7 0 1




page frames





reference string

7 0 1 2 0 3 0 4 2 3 0 3 2 1 2 0 1 7 0 1



page frames





reference string

7 0 1 2 0 3 0 4 2 3 0 3 2 1 2 0 1 7 0 1

7	7	7
	0	0
		1

page frames





reference string



7 0 1 2 0 3 0 4 2 3 0 3 2 1 2 0 1 7 0 1

7	7	7	2
	0	0	0
		1	1

page frames





reference string

7 0 1 2 0 3 0 4 2 3 0 3 2 1 2 0 1 7 0 1



7	7	7	2
	0	0	0
		1	1

page frames





reference string

7 0 1 2 0 3 0 4 2 3 0 3 2 1 2 0 1 7 0 1



7	7	7	2	2
	0	0	0	3
		1	1	1

page frames





reference string

7 0 1 2 0 3 0 4 2 3 0 3 2 1 2 0 1 7 0 1



7	7	7	2	2	2
	0	0	0	3	3
		1	1	1	0

page frames





reference string

7 0 1 2 0 3 0 4 2 3 0 3 2 1 2 0 1 7 0 1



7	7	7	2			2	2	4
	0	0	0			3	3	3
		1	1			1	0	0

page frames





reference string

7 0 1 2 0 3 0 4 2 3 0 3 2 1 2 0 1 7 0 1



7	7	7	2			2	2	4	4										
	0	0	0			3	3	3	2										
		1	1			1	0	0	0										

page frames





reference string

7 0 1 2 0 3 0 4 2 3 0 3 2 1 2 0 1 7 0 1



7	7	7	2																
	0	0	0																
		1	1																

page frames





reference string

7 0 1 2 0 3 0 4 2 3 0 3 2 1 2 0 1 7 0 1



7	7	7	2																
	0	0	0																
		1	1																

page frames





reference string

7 0 1 2 0 3 0 4 2 3 0 3 2 1 2 0 1 7 0 1



7	7	7	2			2	2	4	4	4	0								
	0	0	0			3	3	3	2	2	2								
		1	1			1	0	0	0	3	3								

page frames





reference string

7 0 1 2 0 3 0 4 2 3 0 3 2 1 2 0 1 7 0 1



7	7	7	2			2	2	4	4	4	0
	0	0	0			3	3	3	2	2	2
		1	1			1	0	0	0	3	3

page frames





reference string

7 0 1 2 0 3 0 4 2 3 0 3 2 1 2 0 1 7 0 1



7	7	7	2			2	2	4	4	4	0			0					
	0	0	0			3	3	3	2	2	2			1					
		1	1			1	0	0	0	3	3			3					

page frames





reference string

7 0 1 2 0 3 0 4 2 3 0 3 2 1 2 0 1 7 0 1



7	7	7	2																
	0	0	0																
		1	1																

page frames





reference string

7 0 1 2 0 3 0 4 2 3 0 3 2 1 2 0 1 7 0 1



7	7	7	2			2	2	4	4	4	0			0	0		
	0	0	0			3	3	3	2	2	2			1	1		
		1	1			1	0	0	0	3	3			3	2		

page frames





reference string

7 0 1 2 0 3 0 4 2 3 0 3 2 1 2 0 1 7 0 1



7	7	7	2			2	2	4	4	4	0				0	0		
	0	0	0			3	3	3	2	2	2				1	1		
		1	1			1	0	0	0	3	3				3	2		

page frames





reference string

7 0 1 2 0 3 0 4 2 3 0 3 2 1 2 0 1 7 0 1



7	7	7	2			2	2	4	4	4	0			0	0		7		
	0	0	0			3	3	3	2	2	2			1	1		1		
		1	1			1	0	0	0	3	3			3	2		2		

page frames





reference string

7 0 1 2 0 3 0 4 2 3 0 3 2 1 2 0 1 7 0 1



7	7	7	2																	
	0	0	0																	
		1	1																	

page frames





reference string

7 0 1 2 0 3 0 4 2 3 0 3 2 1 2 0 1 7 0 1




7	7	7	2																
	0	0	0																
		1	1																

page frames





FIFO Illustrating Belady's Anomaly





Optimal Algorithm

- Replace page that will not be used for longest period of time
 - 9 is optimal for the example
- How do you know this?
 - Can't read the future
- Used for measuring how well your algorithm performs

reference string

7 0 1 2 0 3 0 4 2 3 0 3 2 1 2 0 1 7 0 1

7	7	7	2		2		2		2		2						7		
	0	0	0		0		4		0		0						0		
		1	1		3		3		3		1						1		

page frames





Least Recently Used (LRU) Algorithm

- Use past knowledge rather than future
- Replace page that has not been used in the most amount of time
- Associate time of last use with each page

reference string

7 0 1 2 0 3 0 4 2 3 0 3 2 1 2 0 1 7 0 1

7	7	7	2		2		4	4	4	0			1		1		1		
	0	0	0		0		0	0	3	3			3		0		0		
		1	1		3		3	2	2	2			2		2		7		

page frames

- 12 faults – better than FIFO but worse than OPT
- Generally good algorithm and frequently used
- But how to implement?





LRU Algorithm (Cont.)

- ❑ Counter implementation
 - ❑ Every page entry has a counter; every time page is referenced through this entry, copy the clock into the counter
 - ❑ When a page needs to be changed, look at the counters to find smallest value
 - ▶ Search through table needed
- ❑ Stack implementation
 - ❑ Keep a stack of page numbers in a double link form:
 - ❑ Page referenced:
 - ▶ move it to the top
 - ▶ requires 6 pointers to be changed
 - ❑ But each update more expensive
 - ❑ No search for replacement
- ❑ LRU and OPT are cases of **stack algorithms** that don't have Belady's Anomaly





Use Of A Stack to Record Most Recent Page References

reference string

4 7 0 7 1 0 1 2 1 2 7 1 2

2
1
0
7
4

stack
before
a

7
2
1
0
4

stack
after
b


↑
a

↑
b






Second-Chance (clock) Page-Replacement Algorithm




(a)

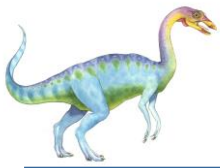


(b)



End of Chapter 9





Paging Example


MAIN MEMORY





Address Translation Scheme

- Address generated by CPU is divided into:
 - **Page number** (p) – used as an index into a **page table** which contains base address of each page in physical memory
 - **Page offset** (d) – combined with base address to define the physical memory address that is sent to the memory unit



- For given logical address space 2^m and page size 2^n





Paging Example


0	a
1	b
2	c
3	d
4	e
5	f
6	g
7	h
8	i
9	j
10	k
11	l
12	m
13	n
14	o
15	p

logical memory

0	5
1	6
2	1
3	2


page table

page number	page offset
p	d
m - n	n



$n=2$ and $m=4$ 32-byte memory and 4-byte pages





CS3600 Operating Systems Concepts

REVIEW





Study Guide

□ The following parts have been covered (textbook: 9th e):

1. Part One: Overview
2. Part Two: Process Management
3. Part Three: Memory Management

□ The following is a study guide that should help you to study and complete this course successfully.

1. Part Two:

a. **Chapter 5: From 5.1 to 5.8**

2. Part Three:

a. **Chapter 8: From 8.1 to 8.5**

b. **Chapter 9: 9.1, 9.2, 9.4**

□ The above shows the recommended sections to study from the textbook. You can remove any sections not covered in class.

□ **Sections in red are the core sections.** You should focus more on these sections.





- Questions will be posted in canvas





Chapter 8: Memory Management

- ❑ Background
 - ❑ Swapping
 - ❑ Contiguous Memory Allocation
 - ❑ Segmentation
 - ❑ Paging
 - ❑ Structure of the Page Table
-
- ❑ Questions will be posted in canvas





Chapter 9: Virtual Memory

- Background
- Demand Paging
- Page Replacement
- Allocation of Frames

□ Questions will be posted in canvas





Extra Points?

Using your own words, write a report comparing Windows vs Linux on CPU Scheduling


Resources:

- Textbook (in 9th e: sec 6.7, p290 to 297)
- Codes are not required
- Do not copy from textbook
- Submission: Online
- Due: Before final exam





Diagram of Process State



Process Management





Process Control Block (PCB)

Information associated with each process

(also called **task control block**)

- ❑ Process state – running, waiting, etc
- ❑ Program counter – location of instruction to next execute
- ❑ CPU registers – contents of all process-centric registers
- ❑ CPU scheduling information- priorities, scheduling queue pointers
- ❑ Memory-management information – memory allocated to the process
- ❑ Accounting information – CPU used, clock time elapsed since start, time limits
- ❑ I/O status information – I/O devices allocated to process, list of open files

Process Management





Threads

- ❑ So far, process has a single thread of execution
- ❑ Consider having multiple program counters per process
 - ❑ Multiple locations can execute at once
 - ▶ Multiple threads of control -> **threads**
- ❑ Must then have storage for thread details, multiple program counters in PCB
- ❑ See next chapter

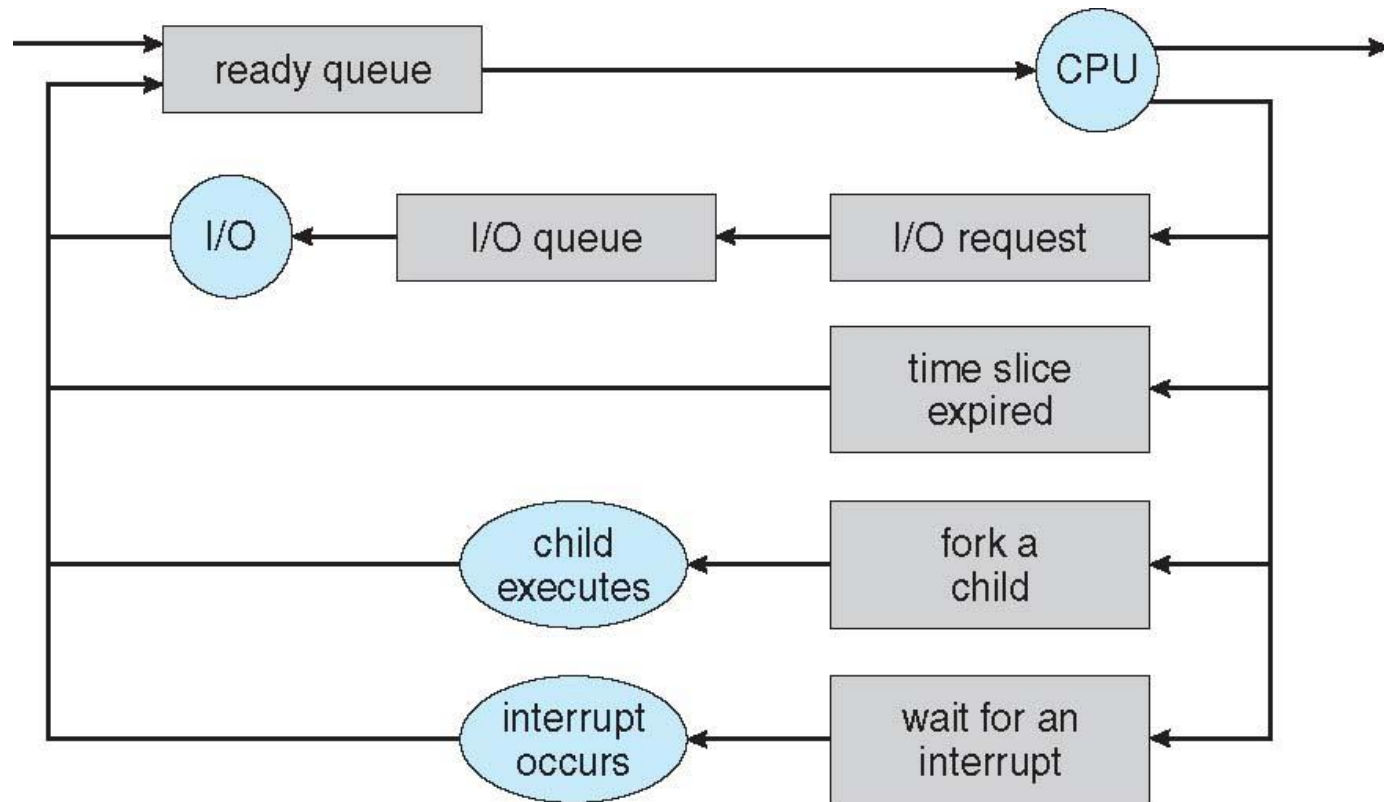
Process vs Threads





Representation of Process Scheduling

- Queueing diagram represents queues, resources, flows

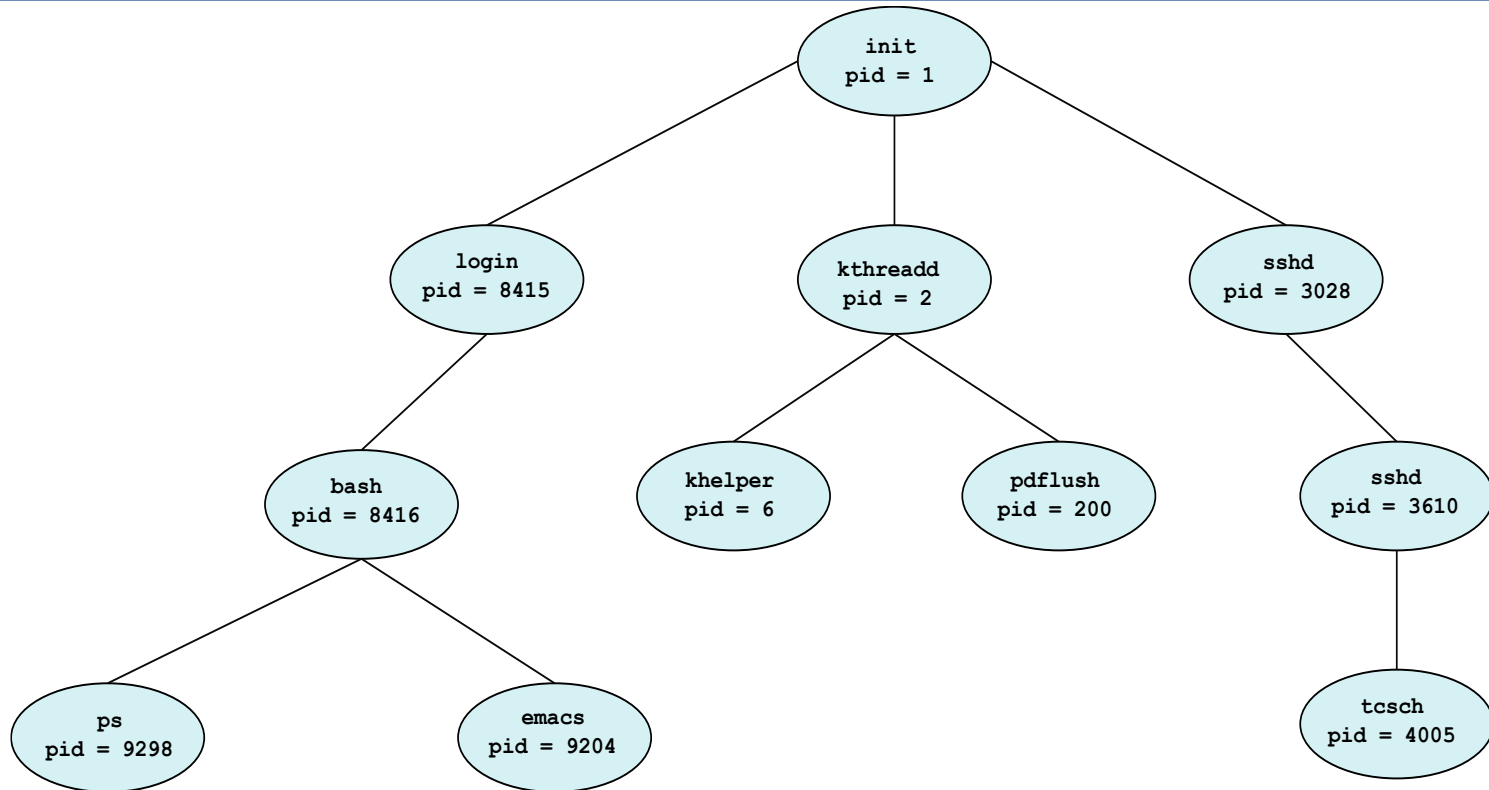


Process Scheduling





A Tree of Processes in Linux



Operations on Process





CPU Scheduler

- Scheduling under 1 and 4 is **nonpreemptive**
- All other scheduling is **preemptive**
 - Consider access to shared data
 - Consider preemption while in kernel mode
 - Consider interrupts occurring during crucial OS activities
- First- Come, First-Served (FCFS) Scheduling
- Shortest-Job-First (SJF) Scheduling
- Priority Scheduling
- Round Robin (RR)

CPU Scheduler






Real-Time CPU Scheduling

- ❑ Can present obvious challenges
- ❑ **Soft real-time systems** – no guarantee when critical real-time process will be scheduled
- ❑ **Hard real-time systems** – task must be serviced by its deadline
- ❑ Two types of latencies affect performance:
 1. Interrupt latency – time from arrival of interrupt to routine that services interrupt
 2. Dispatch latency – time for scheduler to select process off CPU and switch to another

Real Time Scheduling on Process





Methods for Handling Deadlocks

- Deadlock **prevention**
 - Work hard to achieve the prevention, then it will never happen.
- Deadlock **avoidance**
 - Accepting the reality! So, let's “just” avoid it! Do not check what causes it!
- **Detection and Recover**
 - Let it happen, detect it and then recover!
- **Ignore** the problem
 - Lazy, do nothing, because deadlock rarely happens, so why add more functionality to the OS.

Deadlocks

Which method is the best?

